

# Pregel: A System for Large-Scale Graph Processing

M Ravi Kumar   Dr. Ch Sobhan Babu

Mar 2023

# Table of Contents

- 1 Motivation
- 2 Related Work
- 3 Pregel
- 4 Computation Model
- 5 Applications
- 6 References

# Motivation

Many practical computing problems concern large graphs (web graph, social networks, transportation network).

Example :

- Shortest Path
- Clustering
- Minimum Cut
- Connected Components

- Creating a custom distributed framework for every new algorithm.
- Existing distributed framework MapReduce :  
Sub-optimal performance and have usability issues.
- Single-computer graph algorithm libraries like NetworkX,BGL:  
It is not scalable on large data.
- Existing parallel graph systems like parallel BGL:  
These are do not handle fault tolerance and other issues.

Need for a scalable distributed solution

- Google come up with solution, distributed graph parallel computation frame work *Pregel*.
- Vertex centric computation (Think like a vertex).
- Inspired by *Valiant's Bulk Synchronous Parallel* model.
- Scalable and Fault-tolerant platform.
- API with flexibility to express arbitrary algorithm.

- Bulk Synchronous Parallel

- Series of synchronous iterations (supersteps).
- Vertex asynchronously executes some user-defined function in parallel in each superstep.

- Message-passing Model

- Vertex reads messages sent in previous superstep.
- Vertex sends messages, to be read by other vertices in the next superstep.
- Vertex updates states of itself and its outgoing edges.

# Bulk Synchronous Parallel





Figure: Bulk Synchronous Parallel Model

Sequence of iterations/superstep.

# Bulk Synchronous Parallel



**Figure:** Bulk Synchronous Parallel Model (Single Superstep)

- **Local Computation:** every participating processor/thread may perform local computations.
- **Communication:** The processes exchange data with other process.
- **Barrier synchronization :** The process wait until all other process complete above two steps.



# Vertex State Machine

Execution stops when all vertices have voted to halt and no vertices have messages.




Figure: Vertex State Machine

# Computation Model

- Input
  - Directed graph
- Pregel Computation
  - Partition the vertices and allocate them to CPU's/Threads.
  - Superstep 0: All vertices active, initialize vertex value, send message to the out going neighbours.
  - Superstep 1..N-1
    - Active vertex receive message from previous step.
    - Compute user defined function and update its value.
    - Sends messages to outgoing vertices.
    - Votes to halt if it has no further work to do.
    - Program terminated if all vertices are inactive.
- Output
  - Set of vertex updated values.

# Maximum Vertex Value Example




Figure: Maximum Vertex Value Example

Dotted lines are messages. Shaded vertices have voted to halt.

# Applications

- PageRank
- Shortest Path
- Bipartite Matching
- Semi-clustering

# PageRank

PageRank is link analysis algorithm to identify importance of a document based on the number of references to it and the importance of the source documents themselves.

$A$  =  $A$  is given page

$T_1 \dots T_n$  = Pages that point to page  $A$  (citations)

$d$  = Damping factor

$C(T)$  = No of outgoing links of page  $T$

$N$  = Total no of pages.

$PR(A)$  = PageRank of  $A$

$$PR(A) = (1 - d)/N + d\left(\frac{PR(T_1)}{C(T_1)} + \dots + \frac{PR(T_n)}{C(T_n)}\right)$$

# PageRank

```
class PageRankVertex
: public Vertex<double, void, double> {
public:
    virtual void Compute(MessageIterator* msgs) {
        if (superstep() >= 1) {
            double sum = 0;
            for (; !msgs->Done(); msgs->Next())
                sum += msgs->Value();
            *MutableValue() =
                0.15 / NumVertices() + 0.85 * sum;
        }

        if (superstep() < 30) {
            const int64 n = GetOutEdgeIterator().size();
            SendMessageToAllNeighbors(GetValue() / n);
        } else {
            VoteToHalt();
        }
    }
};
```

Pregel system uses the master/worker model

- Master

- Partition the graph and assign input to workers.
- Keep track of which worker holds which portion.
- Recovers faults of workers.

- Worker

- Load its portion of graph into memory.
- Receive messages from neighboring vertices, process the task.
- Update states of vertices, edges.

- Checkpointing
  - The master periodically instructs the workers to save the state of their partitions to persistent storage system.  
e.g., Vertex values, edge values, incoming messages.
- Failure detection
  - Using regular “ping” messages
- Recovery
  - The master reassigns graph partitions to the currently available workers.
  - The workers all reload their partition state from most recent available checkpoint.



# References

- [1] G. Malewicz, M.H. Austern, A.J. Bik, J.C. Dehnert, I. Horn, N. Leiser, G. Czajkowski, Pregel: a system for large-scale graph processing, in SIGMOD (2010).
- [2] Richard Miller, A Library for Bulk-Synchronous Parallel Programming. in Proc. British Computer Society Parallel Processing Specialist Group Workshop on General Purpose Parallel Computing, 1993.
- [3] Luiz Barroso, Jeffrey Dean, and Urs Hoelzle, Web search for a planet: The Google Cluster Architecture. IEEE Micro 23(2), 2003, 22–28.
- [4] Andrew Lumsdaine, Douglas Gregor, Bruce Hendrickson, and Jonathan W. Berry, Challenges in Parallel Graph Processing. Parallel Processing Letters 17, 2007, 5-20.

# Thank you!